

BSc - Data Mining 1

Topic 05 : Exploratory Data Analysis2

Part 01 : EDA Pass2 3

Dr Bernard Butler

Department of Computing and Mathematics, WIT.
(bernard.butler@waltoninstitute.ie)

Autumn Semester, 2021

Outline

- EDA Process
- Datasets = Tips, Titanic and Algae Blooms
- Identifying and resolving issues (missing value, outliers)
- Generating ToDo list for Feature Engineering/Transformation/Selection

Data Mining (Week 5)

Introduction

Motivating Example

Preparation

Data Handling

Exploring Data 1

Exploring Data 2

Building Models

Prediction

Regression
1

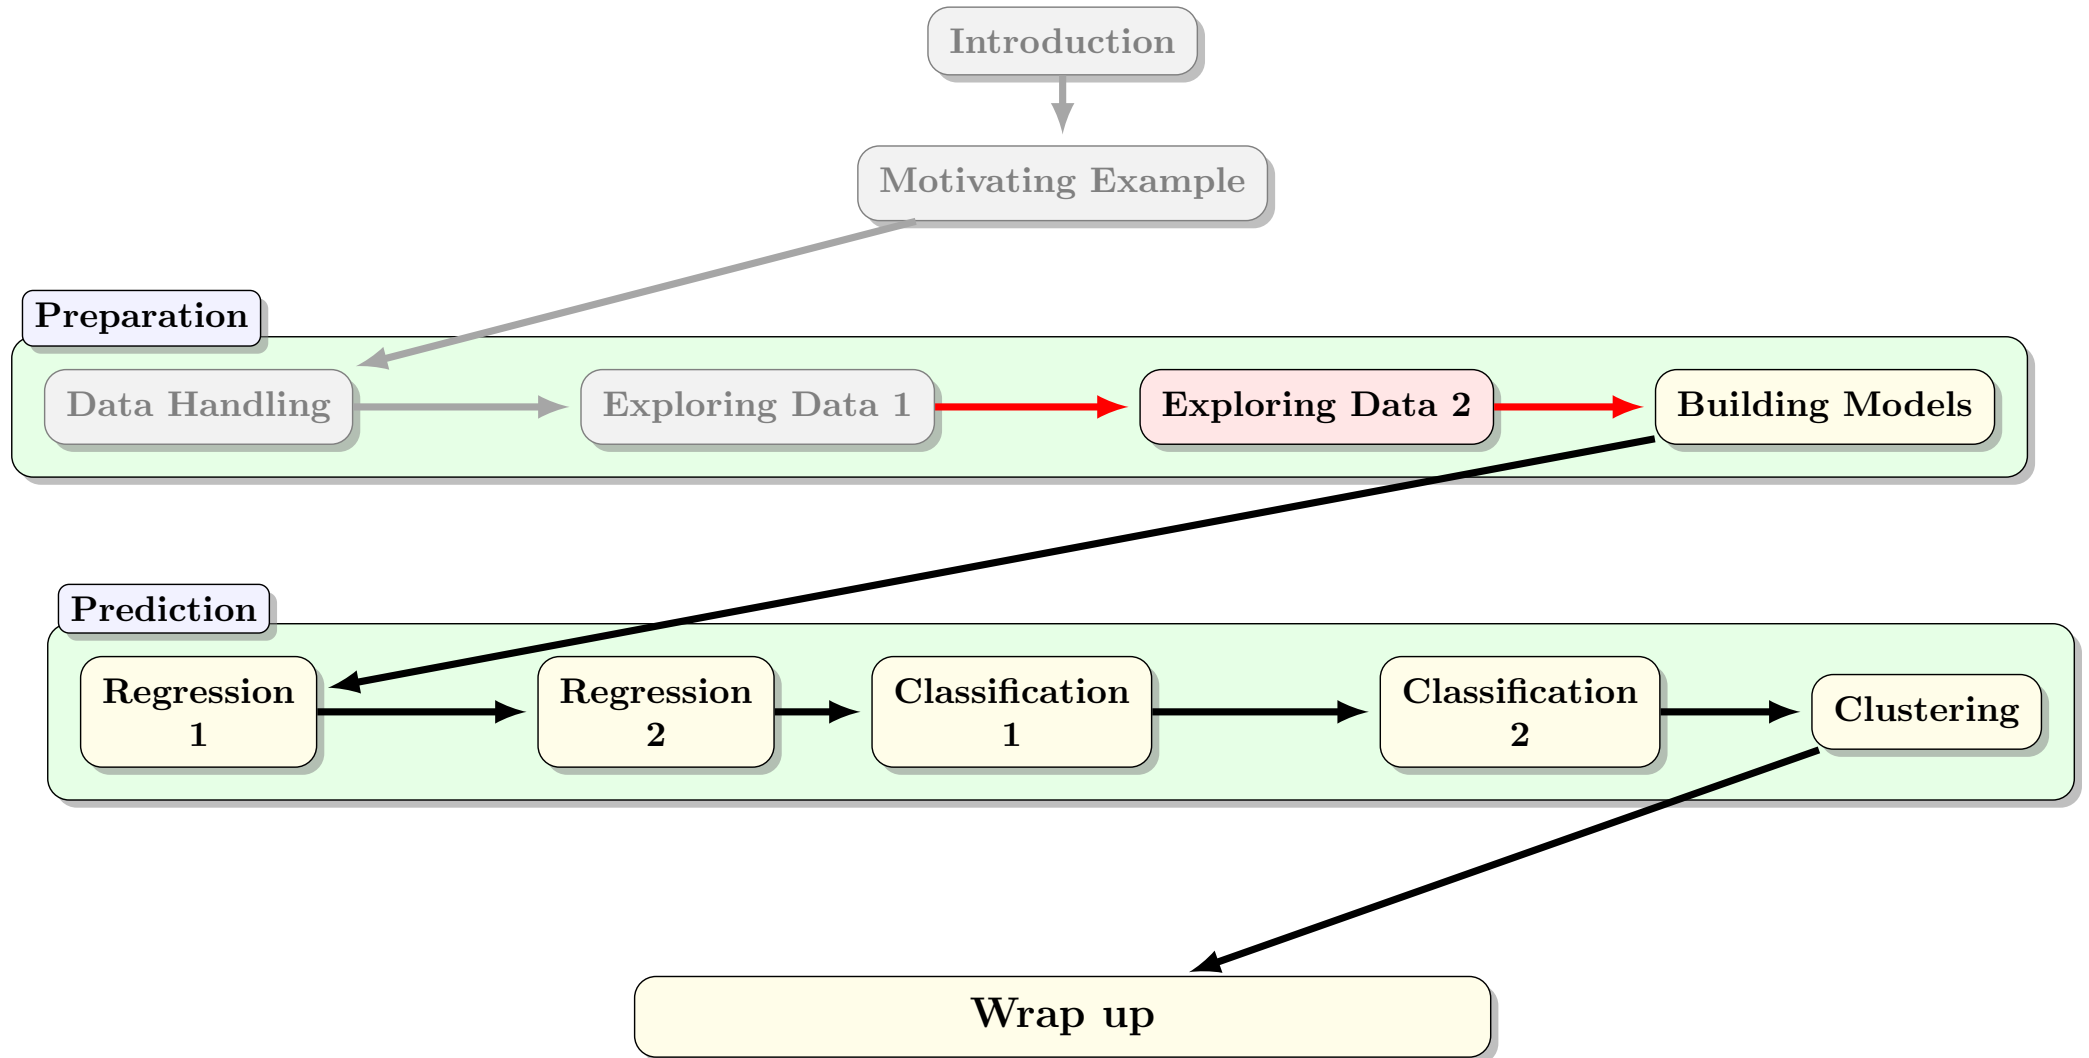
Regression
2

Classification
1

Classification
2

Clustering

Wrap up



EDA Pass2 3 — Summary

1. Review of previous week
2. Second Pass — Individual Features and Target
 - 2.1 Target
 - 2.2 Individual Features
3. Third Pass — Relationships Between Features and Target
 - 3.1 Correlations
 - 3.2 Multi-relation Plots
4. Visualisation - selection of seaborn plots
5. Resources

First Pass — Load Dataset and Initial Clean

- Load dataset
 - Typically either csv or more general “table” format
 - Can be local file or url (read over network)
- Check variables names
 - Should be meaningful and distinct
 - Avoid clashes with reserved words (python or statistical)
- Verify variable types
 - Convert strings to categories, possibly grouping, where possible
 - Ensure numeric data is stored as number (watch out for “Unknown” etc.)
- Identify (and possibly address) missing values
 - Missing values by row or column
 - Leave blank, impute value, drop row/column?

Categorical variables

A Selection of Statistical Visualisations and Metrics

		Relationship to Target	
		Categorical	Numerical
Categorical	Feature		
	nunique, unique, describe, value_counts, ...	crosstab, ...	boxplot, ...
	countplot, ...	countplot, ...	catplot, boxplot, ...
Numerical	describe, ...	groupby+describe, ...	correlations, ...
	histplot, boxplot, displot, qqplot, ...	catplot, boxplot, ...	lmpplot, ...


Numerical Variables

Things here are more complicated as a numerical variable could follow many different distributions. Here we look at data following the standard normal distribution. To start we generate 10,000 values and put in to new DataFrame, df2.

```
rv = stats.norm()
data = rv.rvs(size=10_000)
df2 = pd.DataFrame(data, columns=["x"])
df2.head(5)
```

	x
0	-0.440553
1	0.507263
2	2.791102
3	0.157782
4	-2.740130

```
sns.histplot(x="x", data=df2, kde=True);
```



```
sns.boxplot(x="x", data=df2);
```

Second Pass — Individual Features and Target

- Categorical vs numerical target
 - Categorical vs numerical features
 - Identify (and possibly address) issues
 - Relationship to target.
- } Is it usable?
- } Is it useful?

Dataset: Titanic, Target: Survived


```
df.Survived.value_counts(normalize=True, dropna=False)
```

```
0    0.616162  
1    0.383838  
Name: Survived, dtype: float64
```

```
df.Survived.describe()
```

```
[0, 1]  
Categories (2, int64): [0, 1]
```

```
sns.countplot(x="Survived", data=df);
```



```
df.Survived.unique()
```

```
count    891  
unique    2  
top       0  
freq     549  
Name: Survived, dtype: int64
```

- Simplest classification problem (two classes) with both classes nearly equal frequency.
- In a **unbalanced** classification problem where the minority class occurs about 20% or lower, models can focus on the majority class.

Dataset: Algae Blooms, Target: a1,..., a7


```
targets = [c for c in df.columns if c[0]=="a"]
targets
```

```
['a1', 'a2', 'a3', 'a4', 'a5', 'a6', 'a7']
```

```
df[targets].describe()
```

	a1	a2	a3	a4	a5	a6	a7
count	198.000000	198.000000	198.000000	198.000000	198.000000	198.000000	198.000000
mean	16.996465	7.470707	4.334343	1.997475	5.115657	6.004545	2.487374
std	21.421713	11.065461	6.976788	4.439205	7.511846	11.711053	5.181536
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.525000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	6.950000	3.000000	1.550000	0.000000	2.000000	0.000000	1.000000
75%	24.800000	11.275000	4.975000	2.400000	7.500000	6.975000	2.400000
max	89.800000	72.600000	42.800000	44.600000	44.400000	77.600000	31.600000


```
plt.figure(figsize=(4,6))
sns.histplot(x="a6", data=df);
```



All distributions are heavily skewed to the right, many with outliers (see next slide). All of the zero measurements are probably due to population levels too low to be measured.

Dataset: Algae Blooms, Target: a_1, \dots, a_7

```
fig, axs = plt.subplots(1, 7, figsize=(24,6))
for k, c in enumerate(targets):
    sns.boxplot(data=df, y=c, color="lightblue", ax=axs[k])
    axs[k].set_xlabel(c)
```



The outliers are likely to be true measurements, but their presence can heavily influence the model training — common strategy is to fit two models (one with the case with target outliers and one without) to assess impact of outliers.

Individual Features




To keep this more manageable we will focus more on the Algae Blooms data set ...

	Season	Size	Speed	max_pH	min_O2	mean_Cl	mean_NO3	mean_NH4	mean_oPO4	mean_PO4	mean_Chlor	a1	a2	a
0	winter	small	medium	8.00	9.8	60.800	6.238	578.00000	105.00000	170.00000	50.000	0.0	0.0	0.0
1	spring	small	medium	8.35	8.0	57.750	1.288	370.00000	428.75000	558.75000	1.300	1.4	7.6	4.8
2	autumn	small	medium	8.10	11.4	40.020	5.330	346.66699	125.66700	187.05701	15.600	3.3	53.6	1.9
3	spring	small	medium	8.07	4.8	77.364	2.302	98.18200	61.18200	138.70000	1.400	3.1	41.0	18.

Sneak preview

- Three categorical variables Season, Size, and Speed.
 - No missing values
 - No high cardinality, and reasonable balanced.
- Eight numerical variables max_pH, ..., mean_Chlor
- Missing values present
- Some variables heavily skewed — might need to transform.
- Possibility of features being interrelated — **multicollinearity** — try **principal component analysis**.

Dataset: Algae Blooms, Feature: Season, Target: a1







```
df.groupby("Season")["a1"].agg(["min", "max", "mean", "count", "std"])
```

	min	max	mean	count	std
Season	\bar{x}		n	σ	
spring	0.0	89.8	16.649057	53	23.093786
summer	0.0	64.2	16.038636	44	17.920798
autumn	0.0	86.6	17.745000	40	21.611203
winter	0.0	81.9	17.498361	61	22.568256

- Countplot shows no issues with feature Season — all levels approximately equally represented.
- Catplots show slightly less spread in a1 for Season="summer" observations.
- No/weak relationship between Season feature and a1 target.

Dataset: Algae Blooms, Feature: Size, Target: a1







```
df.groupby("Size")["a1"].agg(["min", "max", "mean", "count", "std"])
```

	min	max	mean	count	std
Size					
small	0.0	89.8	27.255714	70	24.895426
medium	0.0	86.6	11.267857	84	17.163124
large	0.0	56.8	11.611364	44	16.556123

- Countplot shows no issues with feature Size.
- Size="small" rivers have higher frequencies of a1 alga ((point) catplot), and observed frequencies for small rivers is much more widespread across the domain of frequencies than for other types of rivers (boxplot).

Dataset: Algae Blooms, Feature: Speed, Target: a1

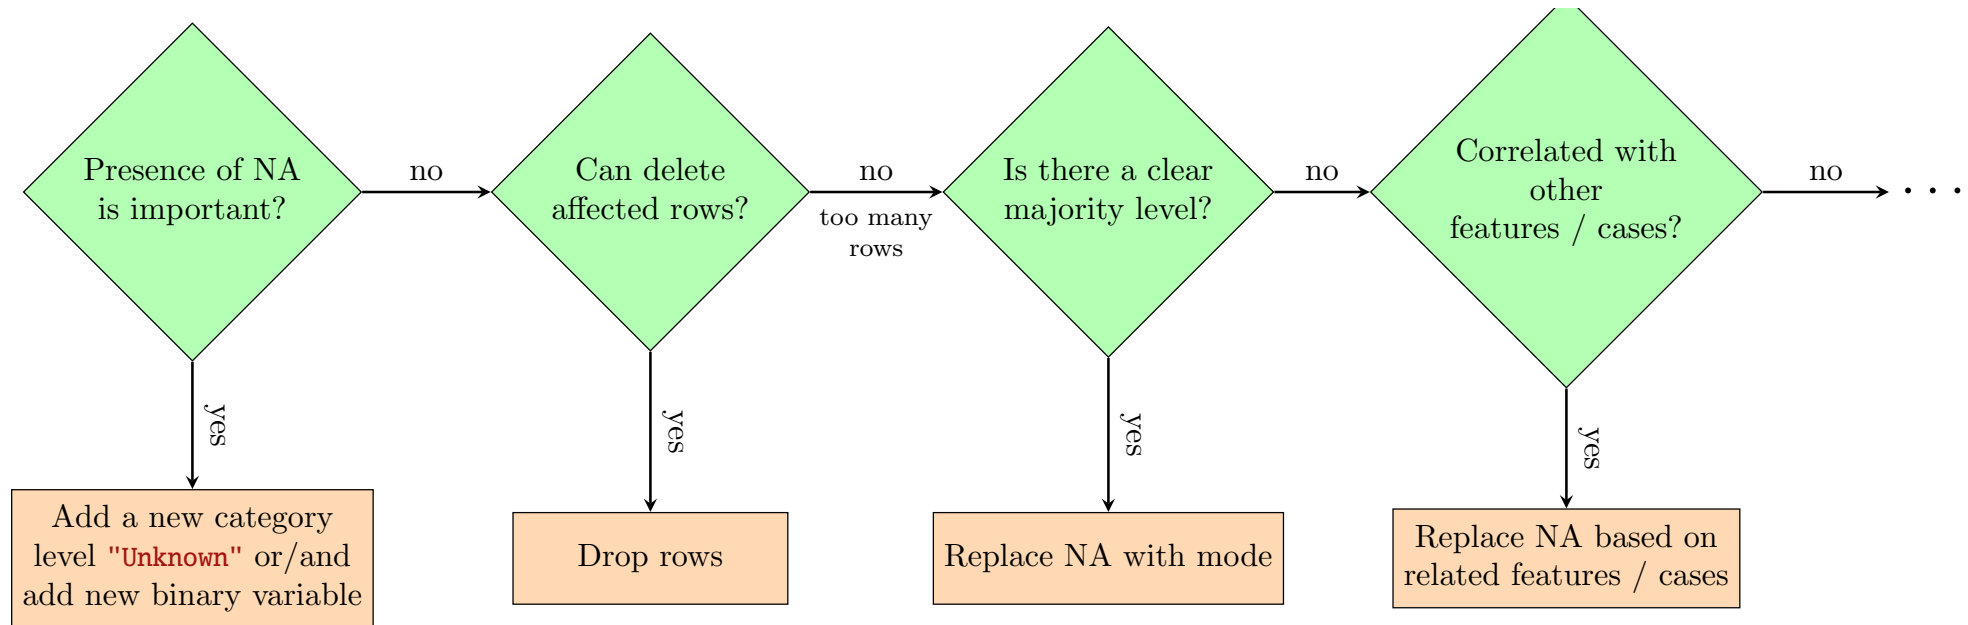


```
df.groupby("Speed")["a1"].agg(["min", "max", "mean", "count", "std"])
```

	min	max	mean	count	std
Speed					
low	0.0	58.7	9.209091	33	13.164758
medium	0.0	89.8	12.548148	81	21.146986
high	0.0	81.9	24.345238	84	22.209123

- Countplot shows no issues with feature Speed.
- Speed="high" rivers have average population of a1 alga ((point catplot), and observed frequencies is much more widespread across the domain of frequencies than for other types of rivers (boxplot).


Categorical Variables — Dealing with Missing Values



In terms of our three datasets, only Titanic has missing values in categorical features:

- Location of cabin's missing values are important (1st class passengers were most likely to have a cabin) so add new category level "Unknown".
- Replace Embarked's 2 missing values with mode ("S", 644/891=72%).
Note: Use `df.Embarked.value_counts(dropna=False)` to include missing values in count tables.



Numerical Variables — Dealing with Missing Values



In terms of our three datasets:

- In Titanic, feature Fare appears to have no missing values, but has 15 zero entries. Are these missing values? or free tickets due to age? ...
- In Algae Blooms, some of the 8 numeric features have NAs ... next few slides.

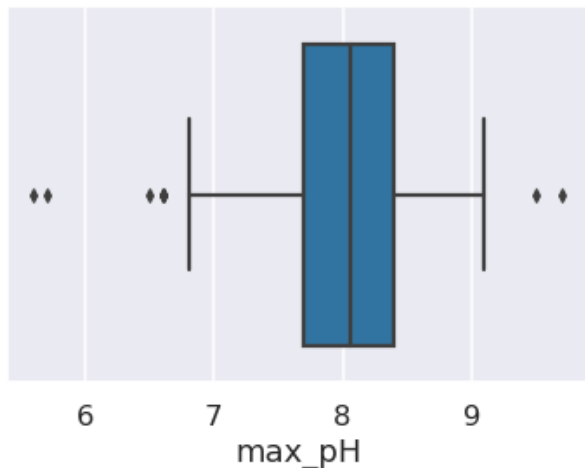
Dataset: Algae Blooms, Feature: max_pH



```
df.max_pH.isna().sum()
```

1

```
count    197.000000
mean      8.019975
std       0.590169
min       5.600000
25%       7.700000
50%       8.060000
75%       8.400000
max       9.700000
Name: max_pH, dtype: float64
```



- Data is relatively normal — minor issue with (left) outliers.
⇒ Will replace (single) NA by mean

```
df.max_pH.fillna(df.max_pH.mean(), inplace=True)
```

Dataset: Algae Blooms, Feature: max_pH, Target: a1

Is there a relationship between feature max_pH and target a1?


```
df[["max_pH", "a1"]].corr()
```

```
sns.lmplot(x="max_pH", y="a1", data=df);
```

	max_pH	a1
max_pH	1.000000	-0.268539
a1	-0.268539	1.000000



(Pearson's) Correlation coefficient, r , measures the strength of a **linear** relationship between two numerical variables.

- near zero means no/weak linear relationship.
- near ± 1 zero means strong linear relationship.
- sign indicates direction.



- Correlation coefficient, $r = -0.27$, shows (at most) a weak negative linear relationship.
- No obvious relationship visible in scatter plot.


Dataset: Algae Blooms, Feature: mean_Chlor



```
df.mean_Chlor.isna().sum()
```

10

count	188.000000
mean	13.971197
std	20.495920
min	0.200000
25%	2.000000
50%	5.475000
75%	18.307500
max	110.456000
Name: mean_Chlor, dtype: float64	



- Data is not normal, heavily skewed to the right \Rightarrow mean is a poor representative of the central location.
 \Rightarrow Will replace (single) NA by median

```
df.mean_Chlor.fillna(df.mean_Chlor.median(), inplace=True)
```

After Target and Individual Feature Pass — Where are we?

Tips

- Reviewed each feature — location, spread, shape, issues.
- No missing values
- total_bill, and total_tip have possible outliers.

Titanic

- Reviewed each feature — location, spread, shape, issues.
- Generated ToDo list for for cleaning, feature extraction
 - Identified features that appear to be related to the target.
 - Feature age has missing values.
 - Feature Fare
 - has 15 measurements with value 0 — decide missing value or not.
 - distribution has large outliers and is skewed — remove/fix outliers and transform.
 - Feature Name has could be used to obtain new feature Title.
 - ...

Algae Blooms

- Reviewed each feature — location, spread, shape, issues.
- Imputed missing values using feature distributions (mean/median).
- Identified features that appear to be related to the target.

Aside: Steps needed to create new feature `Title` from feature `Name`

```
df = pd.read_csv('assets/train.csv')
df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.25	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.28	NaN	S
				Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.92	NaN	S
				Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
				Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

1 Have a (text) feature representing the passenger's name.

Title	Count
Capt	1
Col	4
Don	1
Dona	1
Dr	8
Jonkheer	1
Lady	1
Major	2
Master	61
Miss	260
Mlle	2
Mme	1
Mr	757
Mrs	197
Ms	2
Rev	8
Sir	1
the Countess	1

2 Rare categories can be merged under a single label – 'Dona', 'Lady', 'the Countess', 'Capt', 'Col', etc. . Duplicate/redundant categories can be merged – the titles 'Mlle', and 'Ms' and can be merged under 'Miss'. 'Mme' can be merged with 'Mrs'.

Capt
Col
Don
Dona
Dr
Jonkheer
Lady
Major
Rev
Sir
the Countess

rare_title

Mlle
Ms

Miss

Mme

→ Mrs

3 New categorical feature with 5 unique values

Master	Miss	Mr	Mrs	rare_title
61	264	757	198	29

Third Pass — Relationships Between Features (and Target)

- Correlations

Correlations — Relationship Between two Variables

Pearson's correlation coefficient, r

is a measure of linear correlation between two variables. Its value lies between -1 and +1, -1 indicating total negative linear correlation, 0 indicating no linear correlation and 1 indicating total positive linear correlation.

Spearman's rank correlation coefficient, ρ

is a measure of monotonic correlation between two variables, and is therefore better in catching nonlinear monotonic correlations than Pearson's r . Its value also lies between -1 and +1, with values near zero indicating no monotonic relation.

Kendall rank correlation coefficient, τ

measures ordinal association between two variables. Its value lies between -1 and +1 with values near zero indicating no relation.

Phi-k, ϕ_k


is a new and practical correlation coefficient that works consistently between categorical, ordinal and interval variables, captures non-linear dependency and reverts to the Pearson correlation coefficient in case of a bivariate normal input distribution. Its value also lies between 0 and +1, with values near zero indicating no relation.

Pearson's Correlation Coefficient — Dataset: Algae Blooms

```
columns = df.columns[:12]
corr = df[columns].corr()
corr
```

```
cmap = sns.diverging_palette(230, 20, as_cmap=True)
sns.heatmap(corr, square=True, vmin=-1, vmax=1, cmap=cmap);
```

	max_pH	min_O2	mean_Cl	mean_NO3	mean_NH4	mean_oPO4	mean_PO4	mean_Chlor	a1
max_pH	1.000000	-0.167981	0.136369	-0.130762	-0.093521	0.158769	0.179885	0.445864	-0.268539
min_O2	-0.167981	1.000000	-0.278333	0.099444	-0.087478	-0.416163	-0.487486	-0.153265	0.285564
mean_Cl	0.136369	-0.278333	1.000000	0.225041	0.071913	0.391054	0.457449	0.149856	-0.371171
mean_NO3	-0.130762	0.099444	0.225041	1.000000	0.721444	0.144588	0.168601	0.139679	-0.241211
mean_NH4	-0.093521	-0.087478	0.071913	0.721444	1.000000	0.227237	0.208180	0.088947	-0.132656
mean_oPO4	0.158769	-0.416163	0.391054	0.144588	0.227237	1.000000	0.914365	0.115621	-0.417358
mean_PO4	0.179885	-0.487486	0.457449	0.168601	0.208180	0.914365	1.000000	0.253621	-0.487023
mean_Chlor	0.445864	-0.153265	0.149856	0.139679	0.088947	0.115621	0.253621	1.000000	-0.277987
a1	-0.268539	0.285564	-0.371171	-0.241211	-0.132656	-0.417358	-0.487023	-0.277987	1.000000




- Categorical variables are not included.
- Suggests best predictors for a1 are mean_PO4, mean_oPO4, and meanCl.
- mean_PO4 and mean_oPO4 are highly correlated (0.91) — could use values of one to estimate missing values of the other.

Spearman's Rank Correlation Coefficient — Dataset: Algae Blooms

```
columns = df.columns[:12]
cor = df[columns].corr(method='spearman')
cor
```

```
cmap = sns.diverging_palette(230, 20, as_cmap=True)
sns.heatmap(corr, square=True, vmin=-1, vmax=1,
```

	max_pH	min_O2	mean_Cl	mean_NO3	mean_NH4	mean_oPO4	mean_PO4	mean_Chlor	a1
max_pH	1.000000	-0.148676	0.159079	-0.145182	0.026160	0.290245	0.214569	0.394813	-0.247787
min_O2	-0.148676	1.000000	-0.405142	0.057610	-0.348226	-0.457805	-0.519786	-0.217714	0.283418
mean_Cl	0.159079	-0.405142	1.000000	0.530374	0.592052	0.670399	0.713479	0.564915	-0.546845
mean_NO3	-0.145182	0.057610	0.530374	1.000000	0.425010	0.432303	0.451272	0.346805	-0.382403
mean_NH4	0.026160	-0.348226	0.592052	0.425010	1.000000	0.603157	0.646690	0.406656	-0.449194
mean_oPO4	0.290245	-0.457805	0.670399	0.432303	0.603157	1.000000	0.914921	0.510930	-0.671019
mean_PO4	0.214569	-0.519786	0.713479	0.451272	0.646690	0.914921	1.000000	0.554167	-0.656670
mean_Chlor	0.394813	-0.217714	0.564915	0.346805	0.406656	0.510930	0.554167	1.000000	-0.537823
a1	-0.247787	0.283418	-0.546845	-0.382403	-0.449194	-0.671019	-0.656670	-0.537823	1.000000



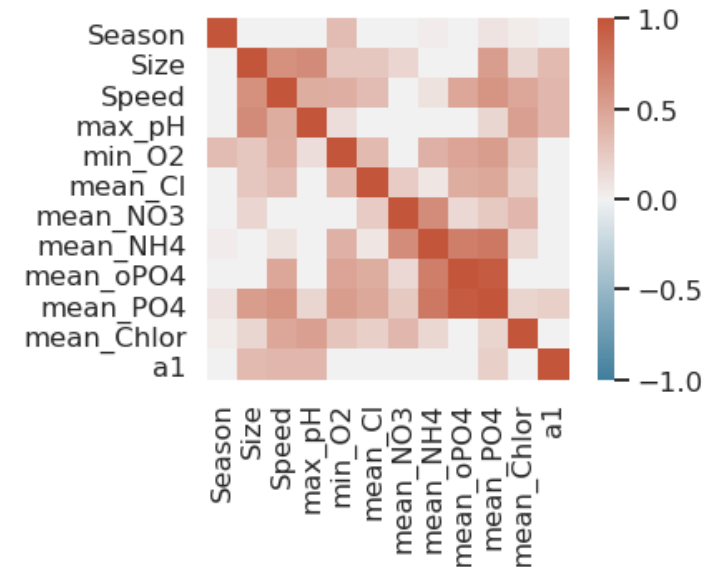
- Now best predictors for a1 also include mean_Chlor and mean_NH4.

Phik Correlation Coefficient — Dataset: Algae Blooms

```
import phik
columns = df.columns[:12]
corr = df[columns].phik_matrix()
corr
```

```
cmap = sns.diverging_palette(230, 20, as_cmap=True)
sns.heatmap(corr, square=True, vmin=-1, vmax=1,
```


	Season	Size	Speed	max_pH	min_O2	mean_Cl	mean_NO3	mean_NH4	mean_oPO4	mean_PO4	mean_C
Season	1.000000	0.000000	0.000000	0.000000	0.343496	0.000000	0.000000	0.034202	0.000000	0.093199	0.04536
Size	0.000000	1.000000	0.620101	0.655207	0.270013	0.268198	0.182410	0.000000	0.000000	0.531635	0.17351
Speed	0.000000	0.620101	1.000000	0.445096	0.437356	0.339237	0.000000	0.101348	0.483298	0.594480	0.47973
max_pH	0.000000	0.655207	0.445096	1.000000	0.125231	0.000000	0.000000	0.000000	0.000000	0.175105	0.52813
min_O2	0.343496	0.270013	0.437356	0.125231	1.000000	0.353196	0.000000	0.416999	0.492457	0.535996	0.29637
mean_Cl	0.000000	0.268198	0.339237	0.000000	0.353196	1.000000	0.243887	0.073692	0.443047	0.472824	0.22558
mean_NO3	0.000000	0.182410	0.000000	0.000000	0.000000	0.243887	1.000000	0.642789	0.158463	0.259915	0.36814
mean_NH4	0.034202	0.000000	0.101348	0.000000	0.416999	0.073692	0.642789	1.000000	0.734681	0.776197	0.16753
mean_oPO4	0.000000	0.000000	0.483298	0.000000	0.492457	0.443047	0.158463	0.734681	1.000000	0.954601	0.00000
mean_PO4	0.093199	0.531635	0.594480	0.175105	0.535996	0.472824	0.259915	0.776197	0.954601	1.000000	0.19292
mean_Chlor	0.045361	0.173516	0.479735	0.528134	0.296376	0.225583	0.368142	0.167533	0.000000	0.192920	1.00000
a1	0.000000	0.353390	0.369374	0.372031	0.000000	0.000000	0.000000	0.000000	0.000000	0.221308	0.00000




- Now include categorical variables — Season is not related, but Size and Speed are.

Multi-Relation Plots


```
sns.boxplot(x="a1", data=df);
```



```
sns.boxplot(x="a1", y="Size", data=df);
```



```
sns.boxplot(x="a1", y="Size", hue="Speed", data=df);
```



- If have n features, then have $n(n - 1)$ possible visualisations — gets a bit crazy.


After Third Pass — Where are we?

- Reviewed each feature — location, spread, shape, issues.
- Identified any correlation among features and with target.
- Located and resolved missing values.
- Generated list of possible feature engineering tasks.

Selected seaborn-based visualisations

- We could easily spend several weeks on EDA visualisation
- There is a long history of visualisation, from infographics to bubble plots
- Seaborn provides a gallery of data science-related visualisation examples
- We consider a selection today that are useful in practice

Histograms with facets



Source: https://seaborn.pydata.org/examples/faceted_histogram.html

What it does


- Facets: show a grid of related plots
- Conditioned by 1 or 2 categorical variables
- Here: flipper length of penguins, by $\text{sex} \times \text{species}$.

When to use it

- Have a key variable, represented by a suitable plot
- Wish to view dependence on 1 or 2 categorical variables in same plot group

Line plots with error bands

Source: https://seaborn.pydata.org/examples/errorband_lineplots.html



What it does


- Multiple numeric variables as lineplots
- Use of colour and linetype
- Overlaid on error bands

When to use it

- Multiple numeric variables on same scale
- Highlight uncertainties

Binning with distribution plots

Source: https://seaborn.pydata.org/examples/hexbin_marginals.html



What it does


- Numeric target with 2 attributes
- Binning provides a heatmap

When to use it

- Numeric target with 2 numeric attributes
- Attributes may be correlated

Contour plots of distributions

Source: https://seaborn.pydata.org/examples/joint_kde.html



What it does


- Penguin bill length \times bill width per species
- Two ways of showing distributions

When to use it

- 2 numeric attributes, split by 1 categorical attribute

Violin plots

Source: https://seaborn.pydata.org/examples/simple_violinplots.html



What it does


- Numeric variable, split by category
- Alternative to boxplot
- data points shown here

When to use it

- Numeric attribute by categorical attribute
- Interested in the shape of the distribution

Scatterplot matrix

Source: https://seaborn.pydata.org/examples/scatterplot_matrix.html



What it does


- Iris data - 4 numeric attributes, 3 species
- All combinations shown

When to use it

- Look at many numeric variables together
- Can use colour or other indicator to show categorical variable

Scatterplot with categorical variables

Source: https://seaborn.pydata.org/examples/scatterplot_categorical.html



What it does


- Show numerical variable in terms of 1 or more categorical variables

When to use it

- More detailed alternative to violinplot

Scatterplot with bubbles

Source: https://seaborn.pydata.org/examples/scatter_bubbles.html



What it does


- Auto mpg data, $\text{mpg} \times \text{horsepower}$
- Plot attributes represent categorical attributes
- Note grouping of numeric variable to create categories

When to use it

- Represent multiple categorical variables in terms of 2 numerical variables

Pointplot for Analysis of Variance

Source: https://seaborn.pydata.org/examples/pointplot_anova.html



What it does


- Trend in pulse rates, by *time* activity (ordered categories)
- Rich plot, with drill down capability

When to use it

- Numeric target as function of multiple categorical attributes

Colour palettes

Source: https://seaborn.pydata.org/examples/palette_choices.html



What it does


- Options for choosing palettes
- Qualitative, Sequential, Diverging

When to use it

- Qualitative: unordered categorical variable
- Sequential: ordered categorical variable
- Diverging: ordered sequential variable

Grouped Violinplots

Source: https://seaborn.pydata.org/examples/grouped_violinplots.html



What it does


- Tips data: total_bill by day × smoker
- Note splits in “violins” to accomodate a category

When to use it

- Adding a second categorical variable to violinplot
- Alternative to faceting

Stacked histograms

Source: https://seaborn.pydata.org/examples/histogram_stacked.html



What it does


- Diamond valuation data
- Show distribution of price by cut
- Be careful: stacked not overlaid!

When to use it

- Can compare histograms by category variable
- Alternative to faceting

Heatmap with scatterplot

Source: https://seaborn.pydata.org/examples/heat_scatter.html



What it does


- Network data
- Highlighting correlated flows
- Use of colour and size of bubbles

When to use it

- Emphasise sign and magnitude of correlations

Heatmap with dendrogram

Source: https://seaborn.pydata.org/examples/structured_heatmap.html



What it does


- Heatmap of correlations
- Dendrogram clusters them to highlight similar values

When to use it


- Need to identify groups of correlated numerical variables

Parallel coordinate and radar plots

Source: https://pandas.pydata.org/docs/reference/api/pandas.plotting.parallel_coordinates.html



Source: <https://www.pythoncharts.com/matplotlib/radar-charts/>



Resources

Resources

Guides

- 1 hour, Youtube on generating seaborn plots — excellent (but wrong on interpretation of box plot)

www.youtube.com/watch?v=6GUZXDef2U0&t=1363s

Articles on Exploratory Data Analysis

- Exploratory Data Analysis (EDA) and Data Visualization with Python

www.kite.com/blog/python/data-analysis-visualization-python/

- When Should You Delete Outliers from a Data Set?

humansofdata.atlan.com/2018/03/when-delete-outliers-dataset

Visualisation

- (Seaborn) Example Gallery

<https://seaborn.pydata.org/examples/index.html>